



USC CSCI 402x Assignment 3

Ted Faber
`faber@isi.edu`

Implement an i-node file system

- Manage disk data structures
 - Free list
 - I-nodes
 - Data blocks
- Implement system calls
 - Read/write
 - Directory ops: link, unlink, new files/devices/directories
- Debugging - kshell

Living in a Real OS



Virtual File System

S5 file system

Memory system (baby VM)

Goal: build a working on-disk file system

- Integrate with VFS
 - Implement specializations for i-node disk-based file system
 - Called s5fs (System V File System)
- Integrate with paging system
 - Implement routines to manage disk-based data structures
- Again, we'll give you acceptance tests
 - Step 1: get the kernel shell up and running

Virtual File System

Vfs File System Object

Vnode Virtual file node

| | |
|--|--|
| | |
| | S5fs Operations read vnode, write vnode, mount, umount |

| | |
|--|--|
| | |
| | S5fs File System Data Free list info I-node info |

| | |
|--|---|
| | |
| | s5-inode Operations Read file, write Link, mknod, |

| | |
|--|---|
| | |
| | s5_inode data Blocks in use Size, etc |

Virtual File System

- Good news:
 - Existing code to make syscalls
 - `kernel/fs/vfs_syscalls.c` (and others)
- Bad news
 - Integration of new code with this somewhat complex system

Virtual File System

- Good news:
 - Existing code to make syscalls
 - `kernel/fs/vfs_syscalls.c` (and others)
 - We give you a valid disk to start
- Bad news
 - Integration of new code with this somewhat complex system

S5 file system superblock

```
/* The contents of the superblock, as stored on disk. */
typedef struct s5_super {
    uint32_t s5s_magic;           /* the magic number */
    uint32_t s5s_free_inode;      /* the free inode pointer */
    uint32_t s5s_nfree;           /* number of blocks currently in
                                   * s5s_free_blocks */

    /** First "node" of free block list */
    uint32_t s5s_free_blocks[S5_NBLKS_PER_FNODE];

    uint32_t s5s_root_inode;      /* root inode */
    uint32_t s5s_num_inodes;      /* number of inodes */
    uint32_t s5s_version;         /* version of this disk format */
} s5_super_t;
```


Where do I get a superblock?

- `s5fs_mount()`
 - In `kernel/fs/s5fs/s5fs.c`
 - Called from the new idle to mount the s5fs
 - Actual disk read isn't working yet
- That code also initializes
 - vfs data structures (including pointer to superblock)
 - vfs function table

S5 inodes

```
/* The contents of an inode, as stored on disk. */
typedef struct s5_inode {
    union {
        uint32_t s5_next_free; /* inode free list ptr */
        uint32_t s5_size;      /* file size */
    } s5_un;
#define s5_next_free s5_un.s5_next_free
#define s5_size      s5_un.s5_size
    uint32_t s5_number;          /* this inode's number */
    uint16_t s5_type;            /* one of S5_TYPE_{FREE,DATA,DIR} */
    int16_t s5_linkcount;        /* link count of this inode */
    uint32_t s5_direct_blocks[S5_NDIRECT_BLOCKS];
    uint32_t s5_indirect_block;
} s5_inode_t;
```


Where do I get an i-node

- V-node routine vget calls s5fs_read_vnode
 - You get to write this
 - Read the i-node from disk and return pointer
- vget is called from s5fs_mount
 - Once you make s5fs_mount work
 - Superblock is read
 - I-node's are readable

Reading and Writing the disk through memory objects

- Manage set of pages (VM-like pages)
 - Pages associated with object
- Operations
 - Fillpage – put data into page (read)
 - Cleanpage – put data on disk (write)
 - Dirtypage - Allocate space for page in fs
- Examples:
 - Disk device implemented
 - V-node – you'll implement s5fs specializations

Pages

- Fixed size memory associated with memory object
 - All data from or going to disk is in a page
 - superblock
 - Data blocks
 - I-nodes (*not* v-nodes!)
- Sent to disk at shutdown or memory pressure
 - All data is write behind

Page frame

/* A pframe structure represents a page frame in physical memory available to the
* kernel. pframes are managed by mmobjs */

```
typedef struct pframe {
```

```
    /* Public read: (do not modify outside pframe.c) */
```

```
    /* Object and page number, which together uniquely identify the page */
```

```
    struct mmobj      *pf_obj;
```

```
    uint32_t          pf_pagenum;
```

```
    /* The address of the page frame. Note that this is NOT a
```

```
    * physical address, but is a virtual address in the kernel's memory
```

```
    * map (i.e., it will be higher than 0xc0000000) */
```

```
    void              *pf_addr;
```

```
    /* Private: */
```

```
    uint8_t           pf_flags; /* PF_DIRTY, PF_BUSY */
```

```
    ktqueue_t         pf_waitq; /* wait on this if page is busy */
```

```
    int               pf_pincount;
```

```
    list_link_t        pf_link; /* link on {free,allocated,pinned}_list */
```

```
    list_link_t        pf_hlink; /* link on hash chain of resident page hash */
```

```
    list_link_t        pf_olink; /* link on object's list of resident pages */
```

```
} pframe_t;
```


Page States

- Busy – moving to/from disk cannot use
- Dirty – memory copy changed
- Pinned – must stay in memory
 - This is different from general pinned description
 - Page may be pinned more than once
 - Keep pin count, unpinned at 0
 - You implement page pinning
 - `kernel/mm/pframe.c`

What To Implement: VM system

- kernel/mm/pframe.c
 - Page reading
 - Page pinning/unpinning
- kernel/fs/s5fs/s5fs.c – vnode ops for paging
 - Make a v-node a memory object
 - s5fs_fillpage, s5fs_cleanpage, s5fs_dirtytypepage

Paging: pframe_get

- Function: retrieve page given memory object and offset
 - Disk: logical disk block
 - v-node: file block
- Logic:
 - Look for copy in memory and return it
 - If low on memory ask pageoutd for more
 - Allocate page and fill it from memory object

Pframe_get: your toolchest

- `pframe_get_resident()` - returns the page if in memory
- `pageoutd_needed()` - true if low on mem
- `pageoutd_wakeup()` - start pageoutd running
 - Wait for it on `alloc_waitq`
- `pframe_alloc()` - get an unused page
- `pframe_fill()` - fill page from memory object
 - Works on disks now
 - You need to make it work for s5fs v-nodes

Page pinning

- 2 lists – alloc_list, pinned_list
 - Use same field in pframe_t so only on one list
- pf_pincount : pinned - unpinned
- pframe_pin()
 - Increment pf_pincount move to pinned_list if not there
- pframe_unpin()
 - Decrement pf_pincount, move to alloc_list if pf_pincount is 0

Making v-nodes into memory objects

- kernel/fs/s5fs/s5fs.c
 - s5fs_fillpage() - put data into page
 - s5fs_cleanpage() - put data onto disk
 - s5fs_dirtytpage() - allocate space
- Note that s5fs v-nodes have their own memory objects too – a disk device

fillpage

- `s5fs_fillpage(vnode_t *vnode, off_t offset, void *pagebuf)`
 - Vnode is the node to read
 - Offset is where to start reading (page aligned)
 - Pagebuf is a page-sized destination
- Find disk block with data (`s5_seek_to_block()`)
- If there is one, copy it out from disk (`read_block`)
- Otherwise copy zeroes out

Why the zeroes?

- Writing an empty file
 - Get page
 - Write data
- Get page will be confused if we don't fill with zeroes...
- When does that page get allocated in the i-node?

dirtypage

- Same signature as fillpage
- If no space allocated in i-node, get some because that page will be written
 - `s5_alloc_block()` - you'll write it
 - Whenever you manipulate i-node use `s5_dirty_inode()` to tell the paging system

cleanpage

- `s5fs_cleanpage(vnode_t *vnode, off_t offset, void *pagebuf)`
 - Vnode is the node to write
 - Offset is where to start writing (page aligned)
 - Pagebuf is a page-sized source
- Find destination disk block (`s5_seek_to_block()`)
- Copy data out to disk (`write_block`)

What to implement: VFS

- vfs operations
 - s5fs_vnode_read (passed in a vnode)
 - Read inode data into page (and pin it)
 - Initialize vnode data structures
 - s5fs_delete_vnode
 - If linkcount is 0 free the i-node
 - Unpin the page

What to implement: v-nodes

- `s5fs_read()`
 - Move data from pages to buffer
- `s5fs_write()`
 - Move data from buffer into pages
 - Allocate pages if needed
 - Adjust size of needed (files only grow)

What to implement: v-nodes

- `s5fs_create()` - make an empty file
 - Allocate an i-node (pull one from disk and pin)
 - Add a directory entry
 - Make modified pages dirty
- `s5fs_mknod()` - Make a device
 - Do all of create and lookup device to link
 - You'll need this for the `kernel_shell`

Directories

```
/* The contents of a directory entry, as stored on disk. */  
typedef struct s5_dirent {  
    uint32_t  s5d_inode;  
    char      s5d_name[S5_NAME_LEN];  
} s5_dirent_t;
```


What to implement: v-nodes

- `s5fs_lookup()` - find a directory entry
 - If name is in the directory, return the vnode associated with it
- `s5fs_link()` - Make a link
 - Add a directory entry linking to the v-node
 - Dirty the file blocks (and maybe the i-node)
- `s5fs_unlink()` - remove a link (not a directory)
 - Remove the directory entry linking to the v-node
 - Reduce the link count
 - Dirty the file blocks (and maybe the i-node)

What to implement: v-nodes

- `s5fs_mkdir()` - make a directory entry
 - Allocate an i-node (pull one from disk and pin)
 - Add a directory entries
 - New directory, `.` and `..`
 - Get the link counts right
 - Make modified pages dirty
- `s5fs_rmdir()` - Unlink an empty directory
 - Remove the directory entry linking to the v-node
 - Reduce the link count
 - Dirty the file blocks (and maybe the i-node)

What to implement: v-nodes

- `s5fs_readdir()` - read a directory entry
 - Read an s5fs directory entry
 - Translate to generic format
- `s5fs_stat()` - file information
 - Read an i-node info
 - Translate to generic format

What to implement: helpers

- All these in `s5fs_subr.c`
- `s5_seek_to_block()` - find and allocate data blocks
 - Translate offset to block containing it
 - If requesting allocation
 - get a data block
 - Connect to i-node
 - Dirty both
 - Used by `fillpage`, `cleanpage`, `dirtypage`

What to implement: helpers

- `s5_write_file()`, `s5_read_file()`
 - Implementations of read and write earlier
 - Useful as subroutines in all directory implementations
- `s5_alloc_block()`
 - Get a new data block
 - Manipulates free list (more in a sec)
 - Dirty new block and superblock

What to implement: helpers

- `s5_find_dirent()`
 - Read directory and match name
 - Return i-node number
- `s5_remove_dirent()`
 - Find entry to remove
 - In not last, write last over it
 - Use `s5_read_file` and `s5_write_file` to dirty blocks
 - Adjust size

What to implement: helpers

- `s5_link()`
 - Add a new directory entry
 - Dirty the files – use read/write
 - Used by all the creation calls

Already implemented: helpers


- `s5_free_block()` - free a data block
- `s5_alloc_inode()` - allocate an i-node
- `s5_free_inode()` - free an i-node and all blocks
- Useful to have – useful to read!

Data Structures: Indirect Blocks

- S5_NDIRECT direct blocks
- After that 1 indirect block
 - A data block that holds S5_NDIRECT pointers
- Take into account in
 - s5_read_file, s5_write_file, s5_seek_to_block

Data Structures: Free List

Superblock
S5_NBLKS_PER_FNODE-1
pointers to free datablocks
Last one points to free data block
With S5_NBLKS_PER_FNODE-1
Pointers to free data blocks.



Data Structures: Free List

- When all pointers in superblock are used
 - Copy the pointers from the “next” free block into superblock
 - Allocate the “next” free block
- `s5_free_block()` does the reverse and is implemented....

Data Structures: Sparse Files

- Try this
 - Create File
 - Seek out 50KB
 - Write a block
- A sparse file will have one block allocated and connected to the i-node
- Make sparse files
 - Think about how `s5fs_dirty` interacts...